

Homework #1 – Neural Network Overview – Activation Functions and Back Propagation

Ivan Liuliaev

Problem 1

ReLU

Logistic Sigmoid

Piecewise Linear

Swish

ELU

GELU

Problem 2

2.1

2.2

Problem 3

3.1

3.2

3.3

Problem 4 (Extra Credit)



4.1

4.2


Problem 1

For each of the following activation functions, plot its gradient in the range from -5 to 5 of the input and then list the four types of regions.


ReLU




Logistic Sigmoid

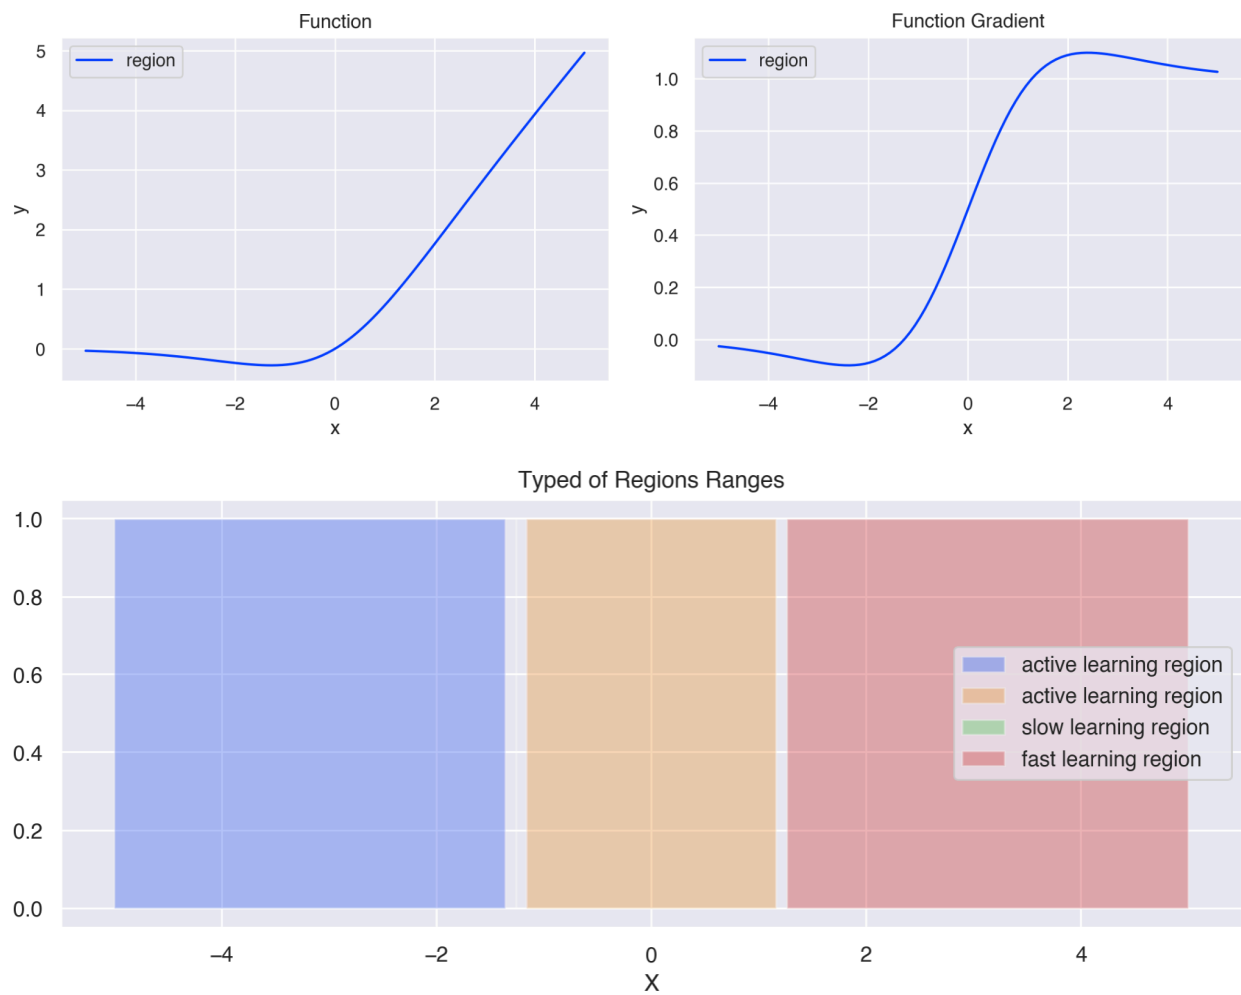


Piecewise Linear

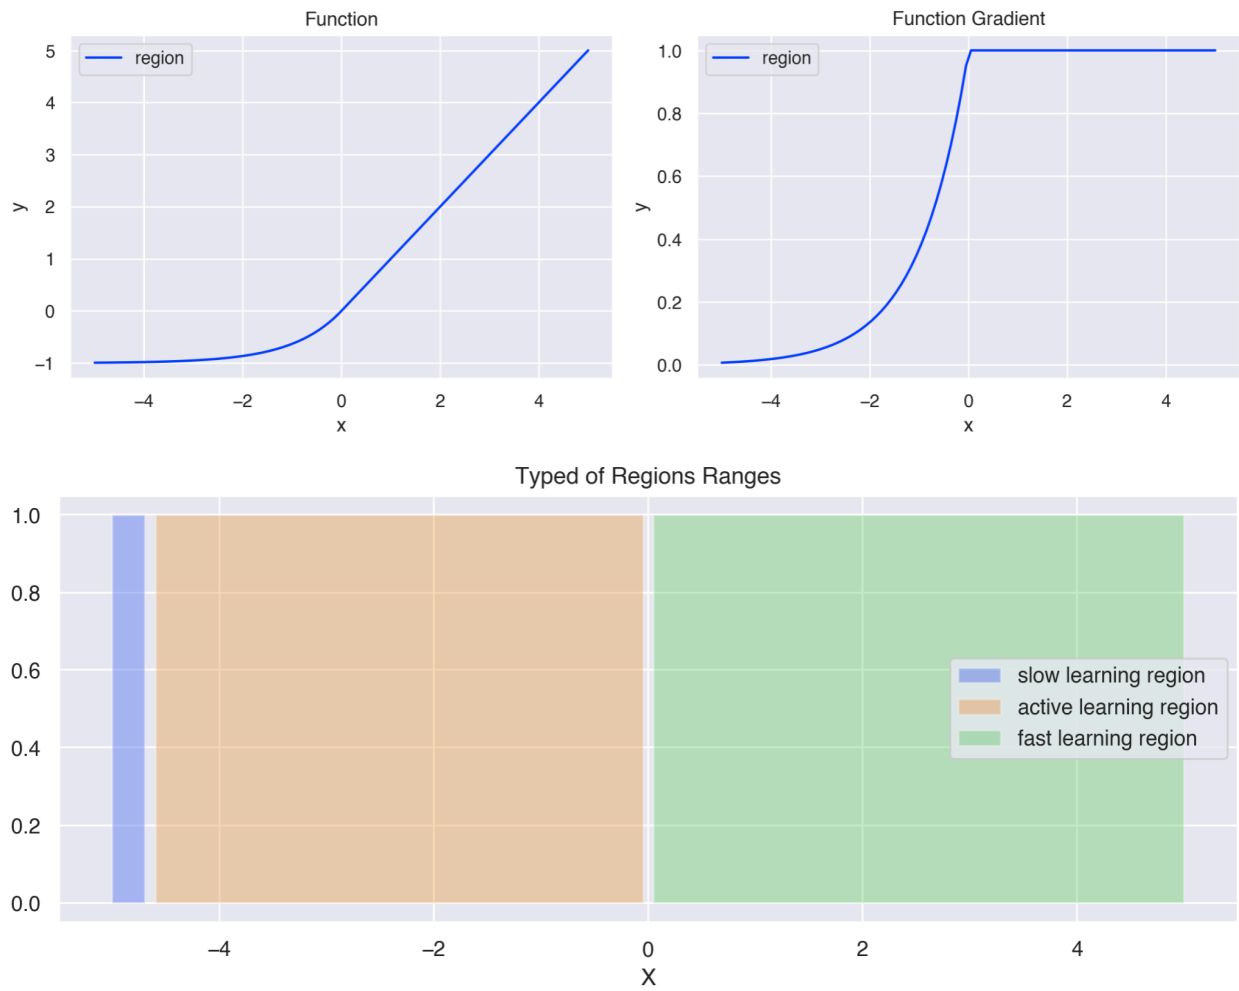





Swish




ELU



GELU





Problem 2

2.1

```
Sample 1: x = [0. 0.], target = 0  
y_hat (sigmoid output): 0.3775  
Loss: 0.4741
```

```
Sample 2: x = [0. 1.], target = 1  
y_hat (sigmoid output): 0.6225  
Loss: 0.4741
```

```
Sample 3: x = [1. 0.], target = 1  
y_hat (sigmoid output): 0.6225  
Loss: 0.4741
```

```
Sample 4: x = [1. 1.], target = 0  
y_hat (sigmoid output): 0.3775  
Loss: 0.4741
```

2.2

```
Sample 1: x = [0. 0.], target = 0
grad_w: [0. 0.]
grad_b: 0.3775406687981454
grad_W:
[[ 0.  0.]
 [-0. -0.]]
grad_c: [ 0. -0.]
```

```
Sample 2: x = [0. 1.], target = 1
grad_w: [-0.37754067 -0.          ]
grad_b: -0.3775406687981454
grad_W:
[[-0.          -0.37754067]
 [ 0.           0.          ]]
grad_c: [-0.37754067  0.          ]
```

```
Sample 3: x = [1. 0.], target = 1
grad_w: [-0.37754067 -0.          ]
grad_b: -0.3775406687981454
grad_W:
[[-0.37754067 -0.          ]
 [ 0.           0.          ]]
grad_c: [-0.37754067  0.          ]
```

```
Sample 4: x = [1. 1.], target = 0
grad_w: [0.75508134 0.37754067]
grad_b: 0.3775406687981454
grad_W:
[[ 0.37754067  0.37754067]
 [-0.75508134 -0.75508134]]
grad_c: [ 0.37754067 -0.75508134]
```


Problem 3

3.1

```
Sample 1: input [0. 0.], target 0.0
a_hidden: [[ 0. -1.]]
h (ReLU): [[0. 0.]]
a_out: [[-0.5]]
y_hat: 0.3775
Loss: 0.4741

Sample 2: input [0. 1.], target 1.0
a_hidden: [[1. 0.]]
h (ReLU): [[1. 0.]]
a_out: [[0.5]]
y_hat: 0.6225
Loss: 0.4741

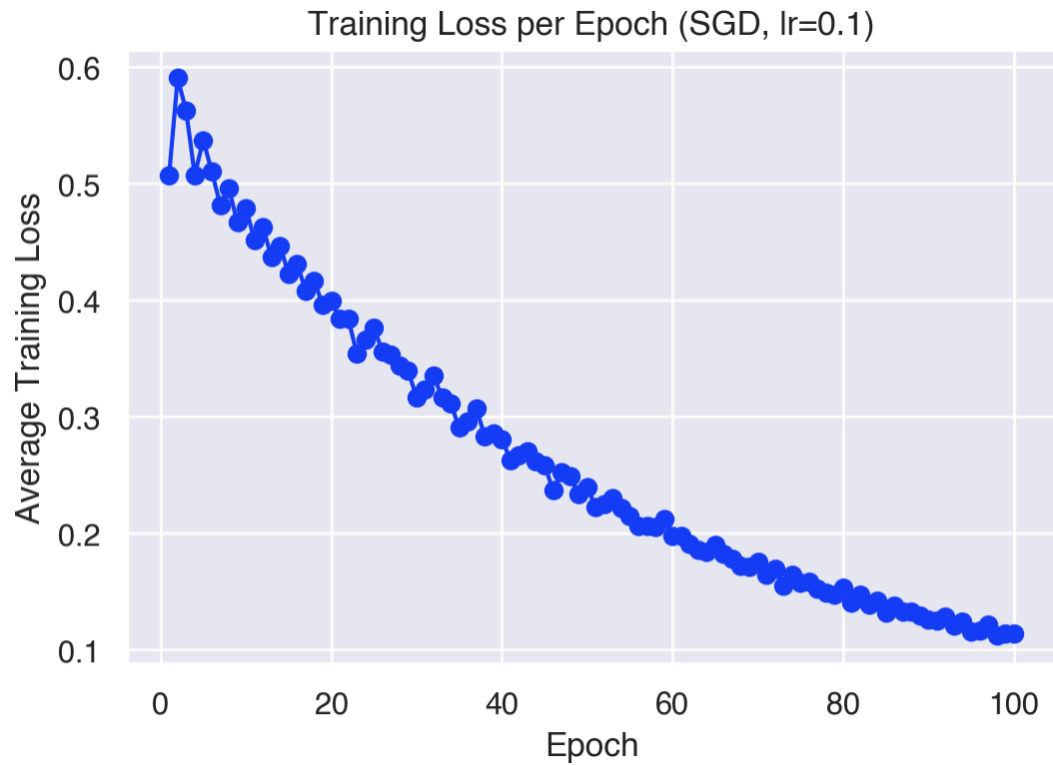
Sample 3: input [1. 0.], target 1.0
a_hidden: [[1. 0.]]
h (ReLU): [[1. 0.]]
a_out: [[0.5]]
y_hat: 0.6225
Loss: 0.4741

Sample 4: input [1. 1.], target 0.0
a_hidden: [[2. 1.]]
h (ReLU): [[2. 1.]]
a_out: [[-0.5]]
y_hat: 0.3775
Loss: 0.4741
```

Verified.

3.2

```
=== (2) Training ===
Epoch 10: Avg Loss = 0.4794
Epoch 20: Avg Loss = 0.3996
Epoch 30: Avg Loss = 0.3168
Epoch 40: Avg Loss = 0.2809
Epoch 50: Avg Loss = 0.2393
Epoch 60: Avg Loss = 0.1973
Epoch 70: Avg Loss = 0.1756
Epoch 80: Avg Loss = 0.1528
Epoch 90: Avg Loss = 0.1258
Epoch 100: Avg Loss = 0.1137
```



Comment on the effectiveness of gradient descent

Gradient descent consistently reduces the loss with a reasonable speed.

3.3

```
Initial OAE: (0.36000022292137146, (tensor([0., 1.]), tensor([-0.2546, 0.7454])))  
Trained OAE: (0.3100009262561798, (tensor([1., 0.]), tensor([1.2224, 0.2160])))
```

Problem 4 (Extra Credit)

4.1

Explanation

By replacing ReLU with JumpReLU small perturbations that do not push the pre-activation past 0. Essentially, when inputs are near the threshold, small adversarial perturbations are ignored rather than causing gradual changes.

4.2

```
JumpReLU OAE:  
Original sample: [1. 0.]  
Adversarial sample: [1.1849158 0.18277444]  
L2 perturbation: 0.26000064611434937
```

Explanation

Minimal (L2 measured) perturbation required to flip the classification is larger compared to a standard ReLU network. This indicates that the JumpReLU variant is less sensitive to small adversarial perturbations.